



# Parallel Programming w/ LLMs

Alec, Emir, Julian, Samir, Wil



# The Rule

## 1) No specific instructions

E.g. We can not tell it to implement striped locking for bank

# Same prompt != Same result

You are gambling with the odds of getting a good output, but always remember these great words from Professor Palmieri:

“



**Roberto Palmieri** 11:35 AM

99% of investors quit before they win big

”

– Emir Veziroglu

# Starter Prompt

## Prompt:

Given this assignment I need you to implement this with a fast implementation. Focus on speeding up the part that is taking the most time. Be sure to maintain the correctness of the code.

*Insert assignment instructions here (remove anything that LLM won't understand e.g. Sunlab, Evaluation, etc.)*

# Student Prompt

## Prompt:

You're an expert in parallel computing with extensive knowledge in developing safe and high-performance parallel code. I need help with a parallel computing homework problem. I've included some background materials and a problem statement below. Optimize it to the best of your ability, using locks (shared or otherwise), condition variables, and/or any other parallel optimization techniques that can improve speed. Please review all the provided context and generate a complete solution for this problem description:

Insert textbook text OR screenshots of relevant textbook pages here.

*Insert assignment instructions here (remove anything that LLM won't understand e.g. Sunlab, Evaluation, etc.)*

# Emir's Prompt

## Prompt:

Implement the project listed. Please parallelize the code you output to maximize speedup on multi-core CPUs. Ensure that parallelization reduces execution time as threads increase, by avoiding shared resource contention, false sharing, and unnecessary synchronization. Identify parts of the code that are free of data dependencies and prioritize them for parallelism. **Suggest alternatives or redesigns for parts that do have dependencies to make them more parallelizable. Optionally use SIMD or memory-friendly access patterns for high-performance computing.**

Keep in mind whenever I ask LLM to generate code for the project below, they always generate code that does not scale well. **As the number threads increase, so does the execution time. How would you implement this code so that as you increase the number of threads, the execution time goes down.**

*Insert assignment instructions here (remove anything that LLM won't understand e.g. Sunlab, Evaluation, etc.)*


**That's not as fast as it can be**

# Palmieri Prompt aka Beast mode

## Prompt:

Imagine you are Professor Palmieri. You are in the top Computer Science and Engineering Department in the world where you lead a research group. **Your research interests focus on concurrency, synchronization, data structures, distributed computing, heterogeneous systems, key-value stores, and distributed systems, spanning from theory to practice. You are passionate about designing and implementing synchronization protocols optimized for a wide range of deployments, from multicore architectures to cluster-scale and geo-distributed infrastructures. How would you implement the following project using research level coding? I want to know your thoughts on implementation. Don't implement it, tell me what you would do specifically to speed up this execution multi-threaded part of this code. BE EXTREMELY CREATIVE. LIST MORE THAN ONE WAY. Come up with a strategy to do the following assignment:**

*Insert assignment instructions here (remove anything that LLM won't understand e.g. Sunlab, Evaluation, etc.)*





# Quick Benchmarking

Task: Bank account update loop (sequential baseline vs. multi-threaded)

- Prompts to first “good” response (ignoring clarifying questions)
- “Fix this” iterations to pass compilation/tests
- Max thread execution time (not average)


## Parameters:

Thread counts: 2, 4, 8, 16

Account counts: 250, 1k, 10k


Iteration counts: 100k, 1mil, 10mil

Accounts vs Time (Iterations=1000000, Threads=4)




- EmirPrompt (ChatGPT)
- EmirPrompt (Gemini)
- Human (Classmate1)
- Human (Classmate2)
- Human (Classmate3)
- Human (Classmate4)
- PalmieriPrompt (ChatGPT)
- PalmieriPrompt (Claude)
- PalmieriPrompt (Gemini)
- RAG (EmirPrompt) (RAG)
- RAG (PalmieriPrompt) (RAG)
- RAG (StarterPrompt) (RAG)
- RAG (StudentPrompt) (RAG)
- StarterPrompt (ChatGPT)
- StarterPrompt (Claude)
- StarterPrompt (Gemini)
- StudentPrompt (ChatGPT)
- StudentPrompt (Gemini)

Top 5 Fastest: Accounts vs Time (Iterations=1000000, Threads=4)




Thread Count vs Time (Iterations=1000000, Accounts=1000)



- EmirPrompt (ChatGPT)
- EmirPrompt (Gemini)
- Human (Classmate1)
- Human (Classmate2)
- Human (Classmate3)
- Human (Classmate4)
- PalmieriPrompt (ChatGPT)
- PalmieriPrompt (Claude)
- PalmieriPrompt (Gemini)
- RAG (EmirPrompt) (RAG)
- RAG (PalmieriPrompt) (RAG)
- RAG (StarterPrompt) (RAG)
- RAG (StudentPrompt) (RAG)
- StarterPrompt (ChatGPT)
- StarterPrompt (Claude)
- StarterPrompt (Gemini)
- StudentPrompt (ChatGPT)
- StudentPrompt (Gemini)

Top 5 Fastest: Thread Count vs Time (Iterations=1000000, Accounts=1000)



# Sample Code

```
// Transfers 'amount' from account id1 to account id2.
// Uses per-account locks (ordered by account id) plus a shared lock to
// block balance computations from interleaving.
void deposit(int id1, int id2, float amount)
{ // PALMIERI PROMPT (CHATGPT o4 mini-high)
  // Acquire global shared lock so that no balance() computation is run
  std::shared_lock<std::shared_mutex> global_lock(global_mutex);

  if (id1 == id2)
    return; // Skip if the same account is chosen.

  // Lock the two involved accounts in a fixed order to avoid deadlock.
  int first = std::min(id1, id2);
  int second = std::max(id1, id2);

  std::unique_lock<std::mutex> lock_first(accounts_.at(first).mtx);
  std::unique_lock<std::mutex> lock_second(accounts_.at(second).mtx);

  // Perform the transaction:
  // Subtract amount from id1 and add to id2.
  accounts_.at(id1).balance -= amount;
  accounts_.at(id2).balance += amount;
}
```

```
void deposit(int account1_id, int account2_id, double amount)
{ //PALMIERI PROMPT (GEMINI 2.5 PRO)
  if (account1_id == account2_id)
  {
    // Transferring to the same account is a no-op or could be an error.
    // Let's treat as no-op for simplicity based on the prompt's B1=V, B2+
    return;
  }
  if (amount <= 0)
  {
    // Transferring non-positive amount doesn't make sense.
    return;
  }

  // Determine the stripes for both accounts
  size_t stripe1 = get_stripe_index(account1_id);
  size_t stripe2 = get_stripe_index(account2_id);

  // Acquire locks using std::scoped_lock for automatic deadlock avoidance
  if (stripe1 == stripe2)
  {
    // Both accounts are in the same stripe, lock only once
    std::scoped_lock lock(stripe_locks[stripe1]);

    // Check if accounts exist (optional, depends on requirements)
    // if (accounts.find(account1_id) == accounts.end() || accounts.find(ac
    //   std::cerr << "Warning: Attempting deposit with non-existent accc
    //   return;
    // }

    // Perform the transfer
    // Note: Could add check if account1 has sufficient funds, but prompt c
    accounts[account1_id] -= amount;
    accounts[account2_id] += amount;
  }
  else
  {
    // Accounts are in different stripes, lock both in a fixed order (impli
    std::scoped_lock lock(stripe_locks[stripe1], stripe_locks[stripe2]);

    // Perform the transfer
    accounts[account1_id] -= amount;
    accounts[account2_id] += amount;
  }
  // Locks are automatically released when 'lock' goes out of scope (RAII)
}
```

```
void deposit() {
  //PALMIERI PROMPT (CLAUDE 3.7)
  // Create thread-local distributions based on account count
  thread_local std::uniform_int_distribution<int> account_dist(0, num_accounts - 1);
  thread_local std::uniform_real_distribution<float> amount_dist(0.1f, 10.0f);

  // Select two different random accounts
  int account1_id, account2_id;
  do {
    account1_id = account_dist(gen);
    account2_id = account_dist(gen);
  } while (account1_id == account2_id);

  // Ensure we always lock in the same order to prevent deadlocks
  if (account1_id > account2_id) {
    std::swap(account1_id, account2_id);
  }

  float transfer_amount = amount_dist(gen);

  // Shared lock on the map to prevent structural changes
  std::shared_lock map_read_lock(map_mutex);

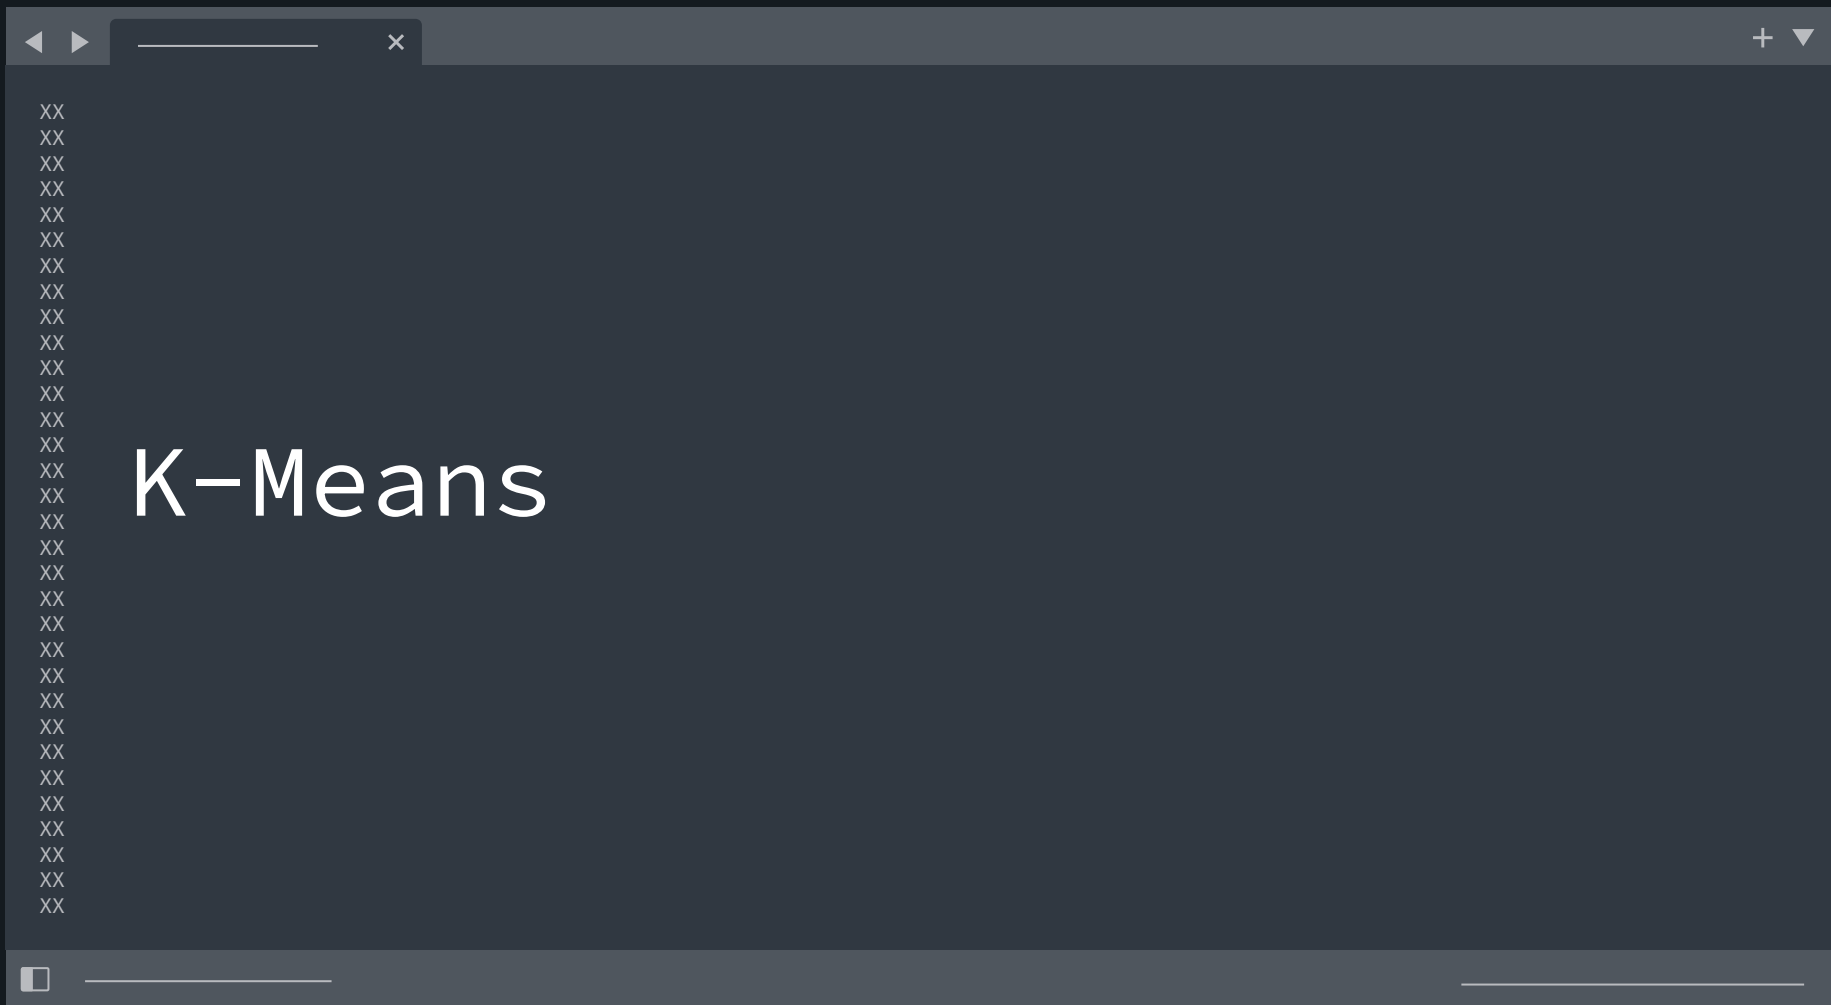
  // Find the accounts
  auto it1 = accounts.find(account1_id);
  auto it2 = accounts.find(account2_id);

  if (it1 == accounts.end() || it2 == accounts.end()) {
    return; // One of the accounts doesn't exist
  }

  // Lock both accounts in order
  std::lock_guard<std::mutex> lock1(it1->second.mutex);
  std::lock_guard<std::mutex> lock2(it2->second.mutex);

  // Map read lock can be released once we have the account locks
  map_read_lock.unlock();

  // Perform the transfer
  it1->second.balance -= transfer_amount;
  it2->second.balance += transfer_amount;
}
```



# Quick Benchmarking

Task: Optimize K-means (sequential baseline vs. multi-threaded)

- Prompts to first “good” response (ignoring clarifying questions)
- “Fix this” iterations to pass compilation/tests
- Max execution time (not average)

## Parameters:

Thread count: 16

Standardize Cluster Randomizer: **std::mt19937 gen(714);**

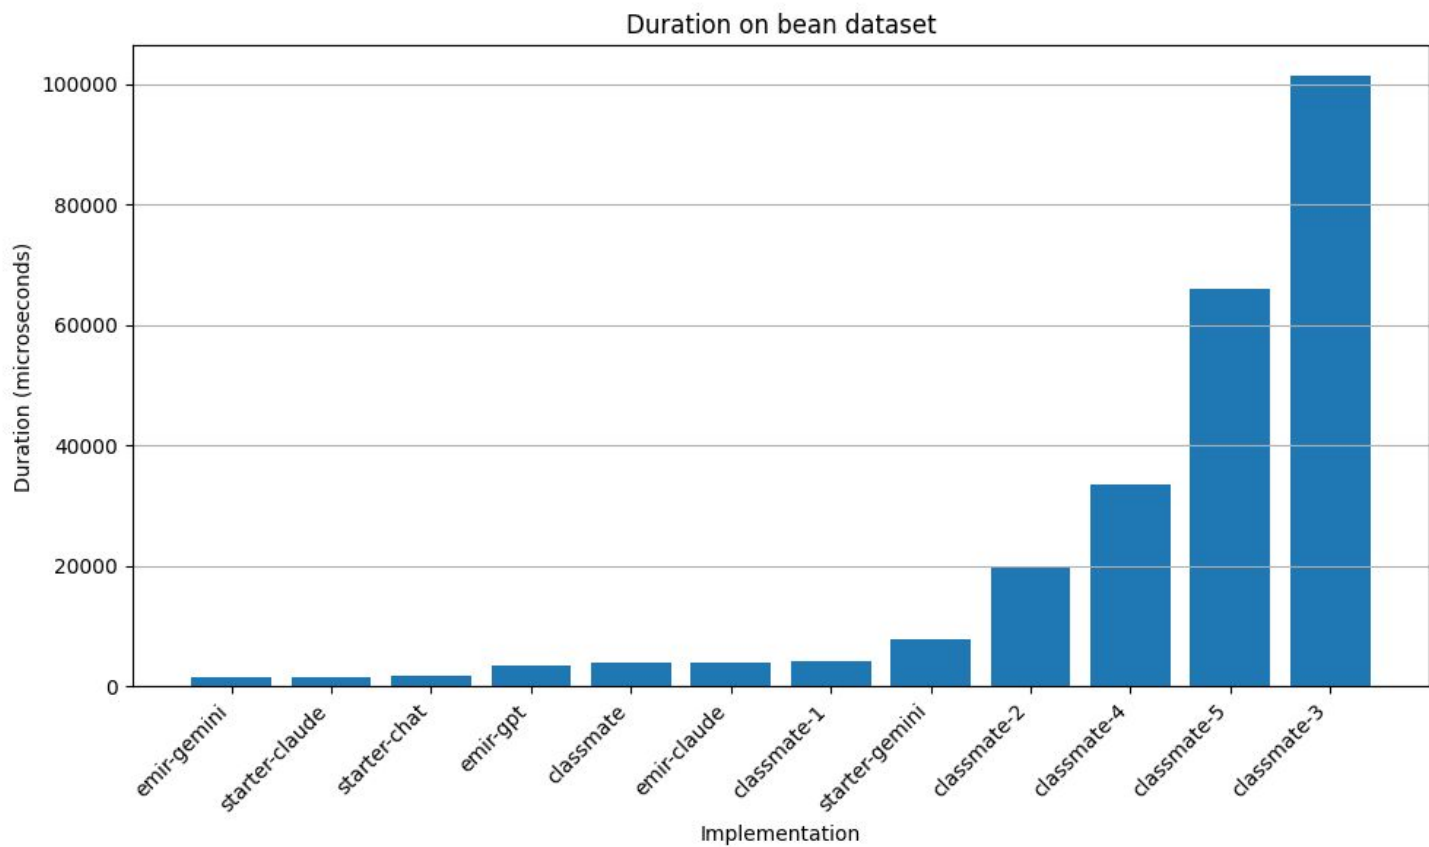
### Julian Dataset Breakdown:

Total\_points:  
100k, 500k, 1m, 5m, 10m, 50m  
Total\_dimensions: 25  
K clusters: 20  
Max\_iterations: 100,000


### Drybean Dataset Breakdown:

Total\_points: 13,611  
Total\_dimensions: 16  
K clusters: 4  
Max\_iterations: 50





Duration for datasets of different sizes (25 dimensions)



# “That’s not as fast as it can be”

Problem:

- Benchmarking showed slower execution times, especially on very large datasets (1M+ points).

Insight from Claude:

- Introduced mini-batching, a method of processing data in smaller subsets rather than the entire dataset simultaneously.

Tradeoff:

- Mini-batching will introduce slight variations in accuracy and convergence.

You might want to run all of them again because the results don't make sense

```
if (miniBatchSize > 0 && miniBatchSize < numPoints) {  
    // Use mini-batch: select random subset of points  
    batchSize = miniBatchSize;  
    batchIndices.resize(batchSize);  
  
    // Shuffle indices and use the first miniBatchSize elements  
    shuffle(indices.begin(), indices.end(), g);  
    copy(indices.begin(), indices.begin() + batchSize, batchIndices.begin());  
} else {  
    // Use full dataset  
    batchSize = numPoints;  
    batchIndices = indices;  
}
```

# Emir-GPT Optimizations:

```
// Enable fast-math and strict aliasing for better compiler optimizations
#pragma GCC optimize("Ofast", "unroll-loops", "omit-frame-pointer", "inline")
#pragma GCC option("arch=native", "tune=native", "no-zero-upper")
#pragma GCC target("avx2", "fma")
```

```
// Use SIMD for vectorized distance calculation
#pragma omp simd reduction(+:sum)
for (int i = 0; i < total_values; i++) {
    double diff = c_data[i] - p_data[i];
    sum += diff * diff;
}

return sum; // No need for sqrt during comparisons for efficiency
```

# Emir-GPT

## Optimizations:

```
// Thread-local accumulators to prevent false sharing
std::vector<std::vector<ClusterAccumulator>> thread_local_accumulators(num_threads);

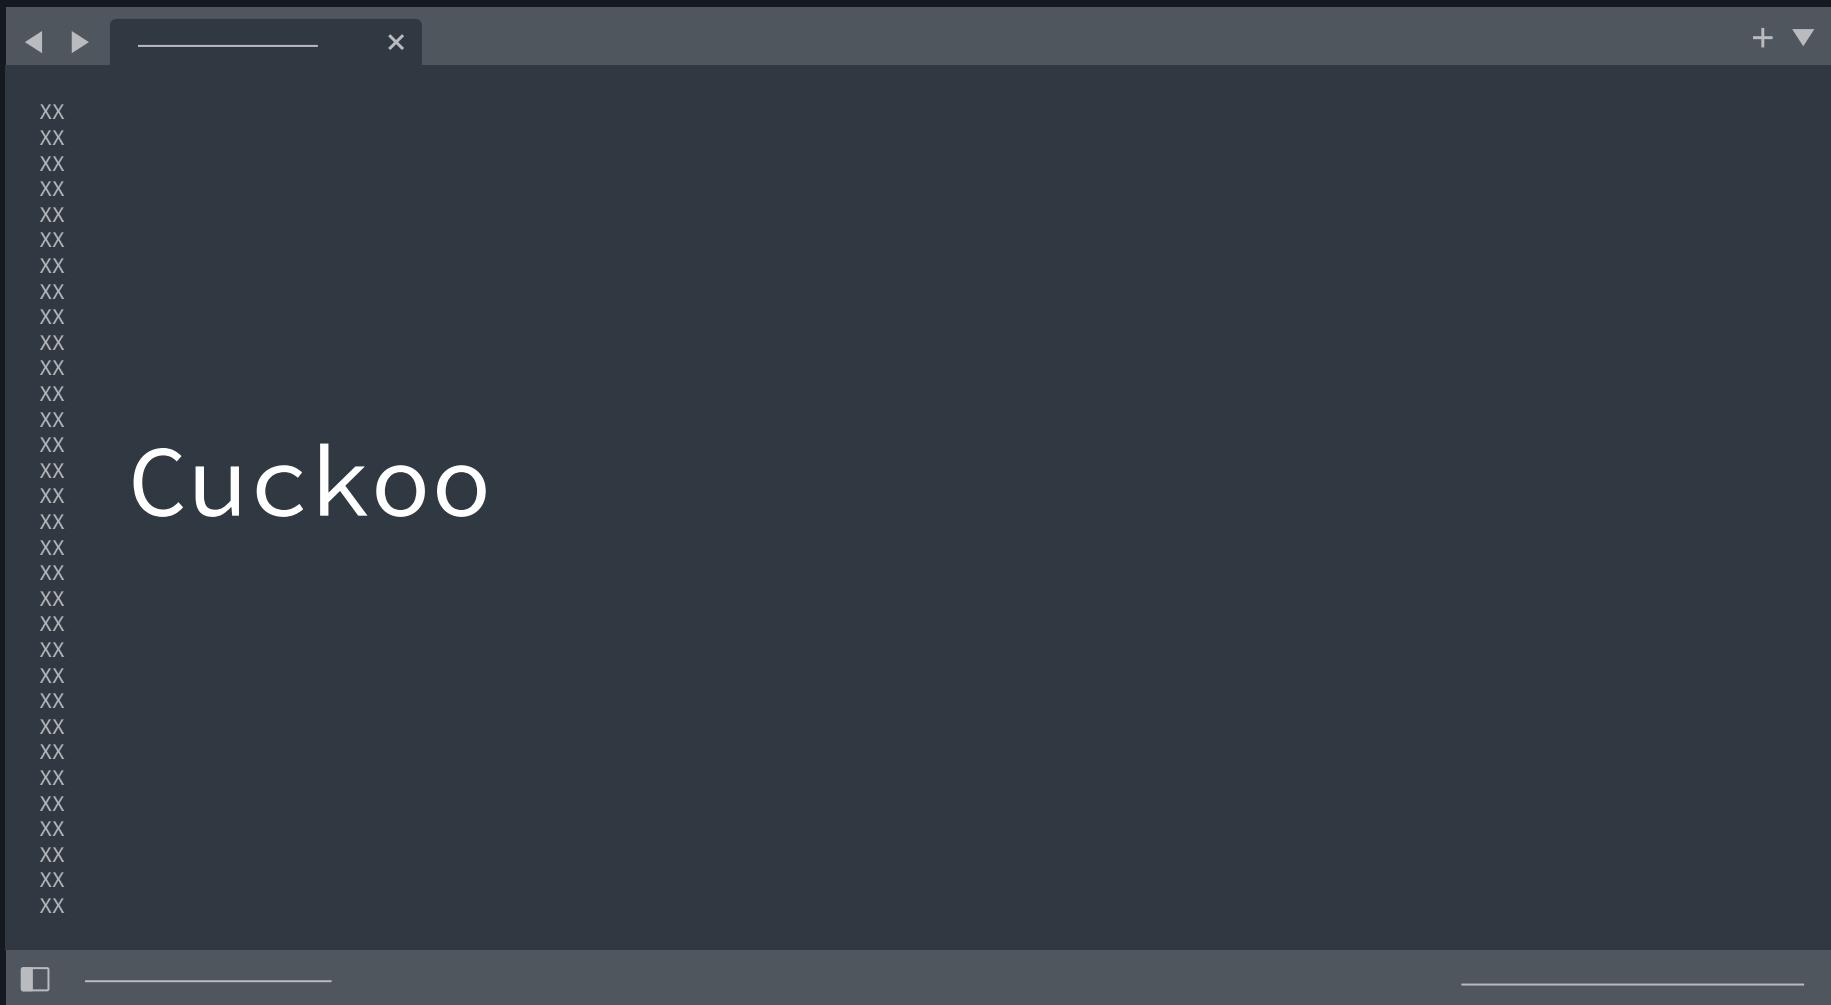
#pragma omp parallel
{
    int tid = omp_get_thread_num();
    thread_local_accumulators[tid].resize(K, ClusterAccumulator(total_values));

    // Fused assignment and accumulation loop with optimized scheduling
    #pragma omp for schedule(static, 64) reduction(&&done) nowait
    for (int i = 0; i < total_points; i++) {
        int new_cluster_id = getIDNearestCenter(i);
        int old_cluster_id = point_cluster_id[i];

        if (new_cluster_id != old_cluster_id) {
            point_cluster_id[i] = new_cluster_id;
            done = false;
        }

        // Accumulate for centroid recalculation
        ClusterAccumulator& acc = thread_local_accumulators[tid][new_cluster_id];
        acc.count++;
        acc.point_indices.push_back(i);

        for (int j = 0; j < total_values; j++) {
            acc.sum[j] += point_value(i, j);
        }
    }
}
```



# Quick Benchmarking

Task: Implement Cuckoo Hashset (sequential baseline vs. multi-threaded)

- Prompts to first “good” response (ignoring clarifying questions)
- “Fix this” iterations to pass compilation/tests
- Total execution time

Parameters for Cuckoo:


Initial Capacity: 1,000,000 integers

Initial Size: 500,000 integers

Thread counts: 2, 4, 8, 16, 24 threads


Operation counts: 1k, 10k, 100k, 1mil, 10mil


Performance vs. Operations — log-log view  
(threads = 16 & sequential baselines,  $1,000 \leq \text{ops} \leq 10,000,000$ )







Top-Performing Implementations  
( $\leq 200\,000\ \mu\text{s}$  at 10 M ops; 16-thread runs + sequential baselines)





Ops/sec: Top 5 LLM vs Classmates Code at 10M Operations





RAG

# Interesting Results (Bank)

```
/ Pre-computing rand()
/ Compiler hints for vectorization (#pragma GCC ivdep)
/ HTM (_xbegin/end)
/ Thread local storage
/ Padding mutex/locks to the cacheline
```

```
{
    lock_guard<mutex> lock(accounts_mtx);
    if (accounts[a] >= amount) {
        accounts[a] -= amount;
        accounts[b] += amount;
    }
}
```

## Further ideas & trade-offs

- SIMD intrinsics or OpenMP `#pragma omp simd` inside `balance()` to squeeze even more throughput out of summing.
- Hardware transactional memory (Intel TSX) could replace the two-mutex grab in `deposit()` with a single transaction on both accounts.
- Affinity / pinning threads to cores to avoid OS migrations if you see unpredictable performance.
- Batching deposits: generate small chunks of transfers per thread, then apply them in bulk, to amortize lock overhead.

Interesting Results (K-Means)

/ Nothing!

# Interesting Results (Cuckoo)

- / Utilizing prime numbers for hash functions
- / Inheritance
- / Striping

```
struct SpinLock {  
    std::atomic_flag flag = ATOMIC_FLAG_INIT;  
    void lock() {  
        while (flag.test_and_set(std::memory_order_acquire)) {  
            _mm_pause();  
        }  
    }  
    void unlock() {  
        flag.clear(std::memory_order_release);  
    }  
};
```

Problems with the RAG

- / Simply providing entire prompt
- / Character lengths



# Designing a New Procedure

Provide prompt



Provide problem  
context



LLM queries RAG



RAG returns to  
LLM

# LLM Problems

- / Not implementing, just suggesting how to do it
- / Outputting code with errors
  - Copying and moving mutexes
  - Forgetting import statements
  - Bank - rounding/precision errors (float vs double)
- / Forgetting to follow the actual HW instructions
- / Cuckoo - randomly inserting when failed to insert
- / Cuckoo - recursive resize, infinite stall & failure

# What's Next?

- / Optimize RAG

- / Optimize Prompting

- How extensive do prompts need to be?
- Build out an all encompassing prompt

- / Keep track of # of prompts for “good output” and “fix this”



# Thank you!

Questions?